

ENTRY POSTURE | PAGE 1

Discover before prescribing.

This plan is a candidate hypothesis, not a claim about Allstate's internal architecture. The first obligation is to understand what current platform choices protect, where teams experience friction, and which evidence should shape the next move.

The goal is not to arrive with a replacement architecture. It is to make the current system legible enough to improve it with the people who own it.

North-star outcome

A governed paved road that is measurably faster and more useful than local reinvention, while making every consequential AI action replayable across identity, context, authority, evidence, reliability, ownership, and economics.

Operating principles

- Tactical empathy before standards.
- Evidence before platform expansion.
- Default paths before exceptions.
- Human authority tied to consequence and reversibility.
- Developer and product experience treated as platform outcomes.
- No target percentage without a baseline and unit.

Stakeholder topology

Executive sponsors; platform and cloud engineering; AI/ML engineering; data platforms; application and product teams; cybersecurity; identity; responsible AI, risk, legal, compliance, and audit; SRE/operations; finance/FinOps; customer, claims, service, agent, and other domain owners; vendors and model providers.

Questions to resolve early

- Where does this role decide, recommend, or convene?
- Which platform services already have trust and adoption?
- Where is the control burden duplicated or delayed?
- Which AI workload offers the best learning value with bounded consequence?
- What must not be disrupted?

Read the system. Define the contract.

WINDOW	OBJECTIVE	ACTIONS	EVIDENCE / DELIVERABLES	DECISION GATE
Days 0-15 Orient	Understand mandate, authority, active priorities, and current platform vocabulary.	Listening sessions; architecture and service-catalog review; incident and exception review; sample workload walkthroughs; identify platform consumers and control owners.	Stakeholder map; decision-rights draft; current vocabulary; evidence inventory; no-regret questions.	Do we agree on the role's decision boundary and the current highest-value unknowns?
Days 16-30 Map	Make current AI demand, paved paths, bypasses, and friction visible.	Map representative workflows from business intent through access, context, orchestration, assurance, production, and outcome; document why local variations exist.	Current-state action traces; friction and duplication map; platform adoption signals; candidate pilot shortlist.	Which friction pattern is reusable enough to solve centrally without erasing legitimate domain variation?
Days 31-45 Frame	Draft the AI Action Contract and risk-sensitive authority model.	Co-design access, work, assurance, and operations planes; define assist/recommend/act tiers; specify minimum trace fields; connect evidence to consequence and reversibility.	Reference-architecture draft; action-trace schema; authority matrix; evaluation and SLO class hypotheses.	Is the proposed contract small enough to adopt, strong enough to trust, and flexible enough for real workflows?
Days 46-60 Align	Resolve conflicts before implementation.	Architecture reviews with engineering, data, security, risk, product, operations, and finance; identify build/buy/reuse choices; test developer journey and exception path.	Decision record; revised reference architecture; pilot charter; baseline measurement plan; known risks and open questions.	Is there one bounded workload, owner, and evidence plan that can prove the contract without a platform-wide rewrite?

Measures established in this window

Current onboarding effort; duplicated implementation; review cycle time; quality and incident evidence available; authority ambiguity; SLO/fallback coverage; cost attribution quality; service adoption; platform-consumer friction. Units are defined with current owners before targets are proposed.

Prove one paved path. Scale the learning.

WINDOW	OBJECTIVE	ACTIONS	EVIDENCE / DELIVERABLES	DECISION GATE
Days 61-75 Instrument	Apply the contract to the selected workload.	Implement or integrate minimum trace fields, authority checks, evaluation evidence, SLO/fallback behavior, ownership, and cost attribution using existing platform capabilities where possible.	Working action trace; test and evaluation evidence; operator/developer feedback; observability and cost data.	Does the paved path improve clarity and evidence without adding disproportionate delivery burden?
Days 76-90 Prove	Run the workload through controlled production or the closest approved evidence environment.	Observe task outcomes, overrides, exceptions, latency, fallback, incidents, review effort, adoption, and cost per completed outcome; conduct a flight-record review.	Pilot evidence review; gap list; architecture and control changes; scale/hold/stop decision.	What is proven reusable, what remains workflow-specific, and what should be stopped or simplified?
Days 91-105 Package	Convert learning into reusable services and guidance.	Publish reference patterns, examples, templates, service-catalog entries, evaluation packs, ownership expectations, and exception pathways; mentor early adopters.	Versioned pattern package; consumer journey; onboarding guide; governance evidence path; operating dashboard.	Can another team adopt the pattern with materially less reinvention and equal or better evidence?
Days 106-120 Prioritize	Establish the next capability roadmap from operating evidence.	Review portfolio demand and platform gaps; prioritize gateway, retrieval, agent orchestration, evaluation, observability, resilience, security, and cost capabilities by repeated need and outcome leverage.	12-month capability roadmap; investment and stop decisions; service ownership; review cadence; mentoring agenda.	Does the roadmap fund reusable constraints and outcomes rather than technology novelty?

Day-120 outputs

- Validated AI Action Contract and reference architecture.
- One evidence-backed paved path and action trace.
- Authority and evaluation patterns by consequence tier.
- Service-catalog candidates and consumer journey.
- Balanced platform scorecard with baselines and owners.
- Prioritized capability roadmap and stop list.

Failure signals to watch

- Platform compliance rises while platform adoption falls.
- Controls remain manual and late.
- Trace volume grows without decision usefulness.
- Cost is measured only at token level.
- Domain teams lose outcome ownership.
- Every exception becomes a new central feature.