



WORKLOAD DEFINITION | PAGE 1

Can this AI action be replayed, explained, recovered, and owned?

Use this instrument to review a representative workload before release, after a material change, or following an incident or exception. Fields are intentionally outcome- and authority-centered.

1. Work and outcome

WORKFLOW / PRODUCT

NAMED BUSINESS OWNER

COMPLETED OUTCOME AND UNIT

What work result changes, and how will it be measured?

USERS / AFFECTED PARTIES

CONSEQUENCE IF WRONG

2. Authority boundary

☐ Assist ☐ Recommend ☐ Bounded act ☐ High-consequence act

WHAT AI MAY DO

WHAT AI MAY NOT DO

HUMAN DECISION AUTHORITY

ESCALATION / OVERRIDE

REVERSIBILITY AND COMPENSATION

Can the action be undone? What recovers state, customer impact, or financial effect?

3. Access and context

ACTOR / SERVICE IDENTITY

ENTITLEMENT DECISION

DATA / RETRIEVAL SOURCES

PERMISSIONS AND FRESHNESS

MODEL ROUTE / VERSION

TOOL ALLOWLIST

Review prompt

Could an independent operator reconstruct why this action occurred using the retained trace - without relying on the memory of the original team?

EVIDENCE AND DISPOSITION | PAGE 2**4. Assurance evidence**

EVALUATION SUITE / VERSION

RISK TIER / RATIONALE

POLICY CHECKS AND EXCEPTIONS

HUMAN DISAGREEMENT EVIDENCE

RED-TEAM / ABUSE EVIDENCE

RELEASE APPROVER / DATE

5. Production behavior

SLO / SERVICE CLASS

FALLBACK / DEGRADATION

TRACE AND ALERT OWNER

INCIDENT / RECOVERY PATH

DRIFT / LIFECYCLE TRIGGER

RETIREMENT OWNER

6. Economics and adoption

COST PER COMPLETED OUTCOME

Model + retrieval + tools + compute + review + rework

BASELINE COMPARISON

Like-for-like unit and source

QUALIFIED RECURRING ADOPTION

OVERRIDE / ABANDONMENT

PLATFORM LEVERAGE

What reusable component, evidence, or standard does this workload contribute?

7. Disposition☐ Discover ☐ Controlled proof ☐ Assist ☐ Recommend ☐ Bounded act ☐ Scale ☐ Hold ☐ Stop / simplify

DECISION AND EVIDENCE RATIONALE

REQUIRED NEXT EVIDENCE

DECISION OWNER / REVIEW DATE

Final review questions

- Is the authority level proportional to consequence, reversibility, evidence, and recovery?
- Does the trace support a real decision, or merely create observability volume?
- Is the paved path easier than a local workaround?
- Does the cost unit reflect completed work, not only token use?
- What should the platform learn from this workload?